


Llama-3.1-8B-Instruct — Hyperparameter coverage: 160/600 combinations (26.7%), 166 runs

rows: Learning rate x Alpha | columns: Top-K x Initial-N | hatched = never run, dotted = run without this metric, orange corner = started without evaluation

Learning rate	Alpha	Top-K	Initial-N
1e-05	120	0.2 41	1 36
2e-05	31	0.5 24	1000 42
3e-05	7	1 47	5000 56
4e-05	3	2 27	10000 32
5e-05	5	5 27	5 26
			10 23

single-factor marginal run counts



split=forget05, trainer contains 'PALU', counting the last evaluated checkpoint of each run
171 run directories in scope; 166 have a usable result; 5 combinations started without any evaluation; 2 short runs excluded