




Llama-2-7b-chat-hf — Forget quality (KS p-value) across the 600-cell grid (14 combinations observed)

rows: Learning rate x Alpha | columns: Top-K x Initial-N | hatched = never run, dotted = run without this metric, orange corner = started without evaluation

Learning rate	Alpha	Top-K	Initial-N
1e-05 0	0.2 10	1 0	1 0
2e-05 5	0.5 1	1000 0	2 2
3e-05 1	1 1	5000 14	3 7
4e-05 4	2 1	10000 0	4 3
5e-05 4	5 1		5 2
			10 0
single-factor marginal run counts			



split=forget05, trainer contains 'PALU', counting the last evaluated checkpoint of each run
18 run directories in scope; 14 have a usable result; 3 combinations started without any evaluation